

AI Engineer Programme

Eight weeks of building, ending in a live technical defence in front of a working practitioner.

Recruiters see a scored project defence, not a list of topics you covered.

THE SPINE — SCOPED BACKWARD FROM THE AI/ML HIRING FUNNEL

01 Solo
Project

02 Solo
Project

03 Pair
Project

04 Capstone
Project

Each project mirrors one layer of a real hiring funnel, so the final defence is not the first time you are asked to explain your own work.

STARTS

1 September 2026

FORMAT

Mentor-led live online sessions

DURATION

8 weeks

SEATS

20 per cohort

FEE

₹9,999 all inclusive



01 · WHY THIS EXISTS

Anyone can finish a course

The rare thing is building something and holding your ground when someone questions it.

Most programmes end the same way: you follow a tutorial, you get a certificate, and the certificate says you attended. Hiring does not work like that. A technical interview at a Chennai product company is a conversation in which someone senior picks up the thing you built and starts asking why — why that retriever, why that chunk size, what happened when it failed, what you would do differently with twice the data.

Candidates do not lose those conversations because they lack knowledge. They lose because they have never had one before. So we built the programme backwards from that conversation and made it the final exam.

02 · THE METHOD

Recruitment-Mirrored Project-Based Learning

RM-PBL is our curriculum spine. Four projects, each one mirroring a layer of a real hiring funnel — so the sequence ends exactly where an interview process ends, with you defending your own work to someone senior.

- **Solo 1** — mirrors the technical phone screen
A small trained model with a real evaluation pipeline: data preparation, training, evaluation. You defend your metric choice, your eval methodology, and one failure mode you found yourself.
- **Solo 2** — mirrors the applied domain round
A RAG pipeline component — document chunking, retrieval, top-k ranking. The defence is followed immediately by an extraction round: cold drills on whichever patterns your own code surfaced.
- **Pair** — mirrors the live ML system design round
A full RAG or agent system built with a partner, then defended at a whiteboard in ten minutes with no notes: data, train, serve, monitor, cost.
- **Capstone** — mirrors the applied domain interview, run as a live project defence
A production-shaped RAG or agent system that you scope and build. The PAT is conducted with your architecture diagram in hand, and the assessor probes every design decision in it.

The DSA thread

Data-structures work is not a separate grind running alongside the projects. Each pattern gets one short, defendable build, wrapped in an AI/ML-shaped artifact — a top-k retrieval ranker for heaps, an embedding cache for hashing, a token-budget allocator for dynamic programming. The pattern and the applied-domain interview reinforce each other instead of competing for your time.

03 · THE AI/ML TRACK

What you actually build

The track is aimed at the work AI teams are hiring for right now — retrieval systems, agents, and the unglamorous engineering that makes them survive contact with real users.

● Retrieval & RAG

Document ingestion, chunking strategy, embeddings, vector stores, and the evaluation loop that tells you whether your retriever is actually working or just confident.

● Agents & tool use

Structured outputs, tool calling, multi-step workflows, and the failure modes that appear when a model is allowed to act rather than only answer.

● Model & data foundations

The ML fundamentals underneath: framing a problem, feature and data quality, training and validation discipline, and reading a result honestly.

● Shipping

Getting a system out of a notebook — an interface, an API, deployment, logging, cost. A project nobody can run is not a portfolio piece.

A detailed week-by-week curriculum, lab manual and capstone catalogue are available on request. Ask on WhatsApp and we will send the current version.

04 · THE DIFFERENTIATOR

The Project Assessment Test

A live oral defence of your own work, scored by a working practitioner who did not teach you.

The PAT is closer to a PhD viva than an exam. You present the system you built, and an assessor who builds for a living spends the session probing it — your design decisions, your trade-offs, the parts that broke. There is no question bank to memorise, because the questions come from your code.

The assessor is deliberately not your mentor. Your mentor knows what you meant. A stranger only knows what you can explain.

PAT SCORECARD — FIVE DIMENSIONS, SCORED 1–5

DIMENSION	WHAT THE ASSESSOR IS TESTING	WEIGHT
Problem framing	Can you explain what the system is for and why it was worth building this way?	20%
Technical depth	Do you understand your own implementation below the level of the library call?	25%
Trade-off reasoning	Can you defend what you chose and articulate what you gave up?	20%
Failure analysis	Do you know where it breaks, why, and what you would do about it?	20%
Communication	Can a senior engineer follow you without already knowing the answer?	15%
Scores are recorded against your portfolio. A weak PAT is not hidden — it is feedback you can act on and re-sit.		

05 · WHAT YOU LEAVE WITH

Evidence, not adjectives

- A verified portfolio**
 Four projects with your reasoning attached, hosted and shareable. Verified means a named practitioner examined the work, not that a system checked a box.
- A scored PAT record**
 Your defence, scored across the five dimensions above. A recruiter can see how you performed, not just that you finished.
- A mentor recommendation letter**
 Written by the person who watched you work. It is earned rather than issued — if the work does not merit it, we will tell you what would.
- A Career Hub listing**
 Your verified profile, visible to recruiters and employers in our network. Free to you.

06 · STRAIGHT ANSWERS

What we do not promise

We do not guarantee placement

Anyone who guarantees you a job is either charging you for it or lying to you. What we do is make you materially easier to hire — real projects, a defensible portfolio, a scored assessment from a practitioner, and introductions where they fit. The interview is still yours to win.

Employers pay us when they hire through the Career Hub. You never do. That is the whole model, and it is why we would rather send a company three candidates who hold up than thirty who do not.

07 · FIT

Who this is for

A GOOD FIT IF YOU

- Are a final-year student or recent graduate aiming at an AI/ML or data role
- Are a working professional moving from IT services, testing, support or analytics into AI engineering
- Can write basic Python and are willing to be uncomfortable for eight weeks
- Can commit real evening and weekend hours — this is not a watch-at-2x programme
- Want to be questioned on your work rather than congratulated for it

A POOR FIT IF YOU

- Want a certificate quickly and do not intend to build
- Are looking for a guaranteed job offer in writing
- Have no programming exposure at all — start with fundamentals first and come back
- Cannot attend live sessions and expect recordings to substitute for the defence
- Want someone else's capstone with your name on it

Not sure which side you fall on? Book a fifteen-minute call and we will tell you honestly, including if the answer is no.

08 · FEE AND TERMS

₹9,999, all in

The first cohorts bundle two tiers that are normally priced separately — the taught programme and the PAT credential — because we want the first twenty portfolios to be strong ones.

INCLUDED	Live mentor-led teaching, all four projects, mentor review, the scored PAT, recommendation letter, verified portfolio and Career Hub listing
PAYMENT	Razorpay · <u>full payment or instalment on request</u>
REFUND	Full refund if you exit within week 1, no carve-outs and no questions. Processed by Razorpay in 5–7 working days
CAREER HUB	Free to you, always. Employers pay us per hire

09 · QUESTIONS WE GET

Frequently asked

I am from a non-CS background. Can I do this?

Yes, if you can already write basic Python — loops, functions, reading a file. Several of the strongest career changers come from mechanical, electronics and commerce backgrounds. What matters more is the hours you can commit.

What if I fail the PAT?

You get the scorecard and specific feedback, and you can re-sit. A failed first attempt that you then defend properly is a better story in an interview than a score you were handed.

Are sessions recorded?

Teaching sessions are recorded so you can revise. Reviews and the PAT itself are live — the entire point is the unscripted conversation.

Is the teaching in English or Tamil?

Sessions run in English with Tamil used freely for explanation when it helps. Your PAT defence is in English, because that is what the interview will be in.

Why only 20 seats?

Because every capstone gets mentor review and every student gets a live defence with a practitioner. That does not scale by adding names to a spreadsheet.

Who is teaching?

Ganesan M, co-founder and CTO — M.Tech in Artificial Intelligence and around seven years training programming, AI and ML. PAT assessors are working practitioners drawn from outside the teaching team.

NEXT STEP

Come and watch us teach before you pay us anything.

FREE WORKSHOP	Build a RAG system, live — 17 August 2026, 7:00 PM IST · skillcred.in/rag
---------------	--

BOOK A CALL	calendly.com/admin-skillcred/seat-confirmation
-------------	---

WHATSAPP	skillcred.in/whatsapp
----------	---

EMAIL	ganesan.m@skillcred.in
-------	--

SkillCred Private Limited, Chennai. Grievance officer: Ganesan M, CTO. Full Terms of Service and Refund & Cancellation Policy at skillcred.in. This prospectus describes Cohort 01 commencing 1 September 2026; content and schedule may be refined before the start date.